

# Data Analysis

## 1. Import libraries and declare paths to processed data

```
In [6]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.ticker import FuncFormatter
import seaborn as sns
import stats
import os
import re

# load processed datasets
bom = pd.read_csv("../data/processed/bom_clean.csv")
imdb_basics = pd.read_csv("../data/processed/imdb_basics_clean.csv")
imdb_ratings = pd.read_csv("../data/processed/imdb_ratings_clean.cs
v")
tn_budgets = pd.read_csv("../data/processed/tn_budgets_clean.csv")
```

## 2. Merge our datasets

```
In [7]: # ensure year in budgets
tn_budgets["year"] = pd.to_datetime(tn_budgets["release_date"], errors="coerce").dt.year

# create worldwide gross col if missing in bom
if "worldwide_gross" not in bom.columns:
    bom["domestic_gross"] = pd.to_numeric(bom["domestic_gross"], errors="coerce")
    bom["foreign_gross"] = pd.to_numeric(bom["foreign_gross"], errors="coerce")
    bom["worldwide_gross"] = bom["domestic_gross"] + bom["foreign_gross"]

# merge imdb basics + ratings
imdb = pd.merge(imdb_basics, imdb_ratings, on="movie_id", how="left")

# merge budgets with imdb (title + year)
merged = pd.merge(
    tn_budgets,
    imdb,
    left_on=["title_clean", "year"],
    right_on=["title_clean", "start_year"],
    how="left"
)

# merge bom gross
merged = pd.merge(
    merged,
    bom[["title_clean", "year", "domestic_gross", "foreign_gross",
        "worldwide_gross"]],
    on=["title_clean", "year"],
    how="left",
    suffixes=("_x", "_y")
)

# clean currency strings & convert to numeric
money_cols = [
    "production_budget",
    "domestic_gross_x", "worldwide_gross_x",    # from budgets (remove the $ sign)
    "domestic_gross_y", "foreign_gross", "worldwide_gross_y"
]

for col in money_cols:
    if col in merged.columns:
        merged[col] = (
            merged[col]
            .astype(str)
            .str.replace(r"\$,]", "", regex=True)
            .replace("nan", pd.NA)
        )
        merged[col] = pd.to_numeric(merged[col], errors="coerce")
```

```
# ratings numeric
if "averagerating" in merged.columns:
    merged["averagerating"] = pd.to_numeric(merged["averagerating"], errors="coerce")

merged.head()
```

Out[7]:

	<b>id</b>	<b>release_date</b>	<b>movie</b>	<b>production_budget</b>	<b>domestic_gross_x</b>	<b>worldwide_gross_x</b>	<b>title</b>
<b>0</b>	1	Dec 18, 2009	Avatar	425000000	760507625	2776345279	
<b>1</b>	2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	410600000	241063875	1045663875	car
<b>2</b>	3	Jun 7, 2019	Dark Phoenix	350000000	42762350	149762350	
<b>3</b>	4	May 1, 2015	Avengers: Age of Ultron	330600000	459005868	1403013963	av
<b>4</b>	5	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	317000000	620181382	1316721747	si ep.

## Check to see if well merged

```
In [8]: merged[['movie', "production_budget", "worldwide_gross_y", "averagerating']].head(10)
```

Out[8]:

	movie	production_budget	worldwide_gross_y	averagerating
0	Avatar	425000000	NaN	NaN
1	Pirates of the Caribbean: On Stranger Tides	410600000	1.045700e+09	6.6
2	Dark Phoenix	350000000	NaN	6.0
3	Avengers: Age of Ultron	330600000	1.405400e+09	7.3
4	Star Wars Ep. VIII: The Last Jedi	317000000	NaN	NaN
5	Star Wars Ep. VII: The Force Awakens	306000000	NaN	NaN
6	Avengers: Infinity War	300000000	NaN	8.5
7	Pirates of the Caribbean: At Worldâs End	300000000	NaN	NaN
8	Justice League	300000000	6.579000e+08	6.5
9	Spectre	300000000	8.807000e+08	6.8

## Check to see how many values are present to be used in our analysis & chart generation

```
In [9]: print("Total rows:", merged.shape[0])  
  
print("\n% available (not NaN):")  
print("production_budget:", merged["production_budget"].notna().mean())  
print("worldwide_gross_y:", merged["worldwide_gross_y"].notna().mean())  
print("domestic_gross_y:", merged["domestic_gross_y"].notna().mean())  
print("averagerating:", merged["averagerating"].notna().mean())
```

Total rows: 5845

```
% available (not NaN):  
production_budget: 1.0  
worldwide_gross_y: 0.19059024807527802  
domestic_gross_y: 0.21864841745081265  
averagerating: 0.2636441402908469
```

### 3. Data Analysis & Chart Generation

#### 3.1 Set global rules, sort & de for chart generation

```
In [10]: # global styling rules & path for saving our images
plt.rcParams.update({
    "figure.dpi": 120,
    "savefig.dpi": 300,
    "font.size": 12,
    "axes.titlesize": 20,
    "axes.labelsize": 14,
    "axes.grid": True,
    "grid.alpha": 0.25,
    "axes.spines.top": False,
    "axes.spines.right": False
})

def dollars_billions(x, pos):
    return f"${x/1e9:.1f}B"

def dollars_millions(x, pos):
    return f"${x/1e6:.0f}M"

# make sure columns have numeric values
for col in ["production_budget", "domestic_gross_y", "foreign_gros
s", "worldwide_gross_y", "averagerating"]:
    if col in merged.columns:
        merged[col] = pd.to_numeric(merged[col], errors="coerce")

# build genre-exploded dataframe
df = merged.copy()
df["genre_list"] = df["genres"].astype(str).str.split(",")
df = df.explode("genre_list")
df["genre_list"] = df["genre_list"].astype(str).str.strip()

# keep only rows with actual genre values
df = df[df["genre_list"].notna() & (df["genre_list"] != "") & (df["g
enre_list"] != "nan")]

# reduce noise in sample genre
MIN_MOVIES_PER_GENRE = 20
genre_counts = df["genre_list"].value_counts()
keep_genres = genre_counts[genre_counts >= MIN_MOVIES_PER_GENRE].ind
ex
df_g = df[df["genre_list"].isin(keep_genres)].copy()

# output directory to save our charts
out_dir = "../visuals/"
```

## 3.2 Average ratings for each genre

```
In [11]: ratings_df = df_g.dropna(subset=["averagerating"]).copy()


# sort genres by median rating
genre_median = ratings_df.groupby("genre_list")["averagerating"].median().sort_values(ascending=False)
genres_sorted = genre_median.index.tolist()

data_box = [ratings_df.loc[ratings_df["genre_list"] == g, "averagerating"].values for g in genres_sorted]

plt.figure(figsize=(16, 6))
plt.boxplot(
    data_box,
    labels=genres_sorted,
    showfliers=True
)
plt.xticks(rotation=45, ha="right")
plt.xlabel("Genre")
plt.ylabel("Average Rating")
plt.title("Audience Ratings by Genre (min 20 movies/genre)")
plt.tight_layout()
plt.savefig(out_dir + "ratings_by_genre.png")
plt.show()
```

/tmp/ipykernel\_125760/1351667370.py:10: MatplotlibDeprecationWarning : The 'labels' parameter of boxplot() has been renamed 'tick\_labels' since Matplotlib 3.9; support for the old name will be dropped in 3.11.


```
plt.boxplot()
```



## 3.3 Worldwide Gross vs IMdb ratings

```
In [12]: sc_df = merged.dropna(subset=["averagerating", "worldwide_gross_y"]).copy()
sc_df = sc_df[(sc_df["averagerating"] > 0) & (sc_df["worldwide_gross_y"] > 0)]


plt.figure(figsize=(12, 6))
plt.scatter(sc_df["averagerating"], sc_df["worldwide_gross_y"], alpha=0.35, s=22)
plt.yscale("log")
plt.xlabel("Average Rating")
plt.ylabel("Worldwide Gross ($) (log scale)")
plt.title("Audience Ratings vs Worldwide Gross")
plt.gca().yaxis.set_major_formatter(FuncFormatter(dollars_millions))
plt.tight_layout()
plt.savefig(out_dir + "ratings_vs_worldwide_gross.png")
plt.show()
```



### 3.4 Budget vs Worldwide gross

```
In [13]: df2 = merged.dropna(subset=["production_budget", "worldwide_gross_y"])
df2 = df2[(df2["production_budget"] > 0) & (df2["worldwide_gross_y"] > 0)]

plt.figure(figsize=(6,4))
plt.scatter(df2["production_budget"], df2["worldwide_gross_y"], alpha=0.4)
plt.xscale("log")
plt.yscale("log")
plt.xlabel("Production Budget (log scale)")
plt.ylabel("Worldwide Gross (log scale)")
plt.title("Production Budget vs Worldwide Gross")
plt.grid(True)
plt.savefig(out_dir + "prod_budget_vs_ww_gross.png")
plt.show()
```




## 3.5 ROI vs Ratings

```
In [14]: df3 = merged.dropna(subset=["production_budget", "worldwide_gross_y", "averagerating"]).copy()
df3 = df3[(df3["production_budget"] > 0) & (df3["worldwide_gross_y"] > 0)]

df3["roi"] = df3["worldwide_gross_y"] / df3["production_budget"]

plt.figure(figsize=(6,4))
plt.scatter(df3["roi"], df3["averagerating"], alpha=0.4)
plt.xscale("log")
plt.xlabel("ROI (log scale)")
plt.ylabel("IMDb Average Rating")
plt.title("ROI vs IMDb Rating")
plt.grid(True)
plt.savefig(out_dir + "roi_vs_ratings_scatter.png")
plt.show()
```




### 3.6 Plot the average ROI for each genre

```
In [15]: roi_df = df_g.dropna(subset=["production_budget", "worldwide_gross_y"]).copy()
roi_df = roi_df[(roi_df["production_budget"] > 0) & (roi_df["worldwide_gross_y"] > 0)]

roi_df["roi"] = roi_df["worldwide_gross_y"] / roi_df["production_budget"]

roi_by_genre = roi_df.groupby("genre_list")["roi"].mean().sort_values(ascending=False).head(15)

plt.figure(figsize=(14, 6))
plt.bar(roi_by_genre.index, roi_by_genre.values)
plt.xticks(rotation=45, ha="right")
plt.xlabel("Genre")
plt.ylabel("ROI (x times budget)")
plt.title("Average ROI by Genre")
plt.tight_layout()
plt.savefig(out_dir + "avg_roi_by_genre.png")
plt.show()
```



### 3.7 Plot the revenue for each genre

```
In [16]: rev_df = df_g.dropna(subset=["domestic_gross_y", "foreign_gross", "worldwide_gross_y"]).copy()
rev_df = rev_df[(rev_df["domestic_gross_y"] >= 0) & (rev_df["foreign_gross"] >= 0) & (rev_df["worldwide_gross_y"] >= 0)]


genre_revenue = rev_df.groupby("genre_list").agg(
    domestic=("domestic_gross_y", "sum"),
    foreign=("foreign_gross", "sum"),
    worldwide=("worldwide_gross_y", "sum")
).sort_values("worldwide", ascending=False)

# take top 'N' genres for readability
TOP_N = 12
genre_revenue = genre_revenue.head(TOP_N)

x = np.arange(len(genre_revenue.index))
dom = genre_revenue["domestic"].values
forg = genre_revenue["foreign"].values # stacked to reach worldwide

plt.figure(figsize=(12, 6))
plt.bar(x, dom, label="Domestic")
plt.bar(x, forg, bottom=dom, label="Foreign")

plt.xticks(x, genre_revenue.index, rotation=45, ha="right")
plt.xlabel("Genre")
plt.ylabel("Revenue ($)")
plt.title(f"Revenue by Genre (Top {TOP_N}) — Stacked Domestic + Foreign")
plt.gca().yaxis.set_major_formatter(FuncFormatter(dollars_billions))
plt.legend()
plt.tight_layout()
plt.savefig(out_dir + "revenue_by_genre.png")
plt.show()
```



### 3.8 Plot average world wide revenue for movies released in each month

```
In [17]: # pick gross columns
gross_pref = "worldwide_gross_y"
grossFallback = "worldwide_gross_x" if "worldwide_gross_x" in merge
d.columns else None

# build base month data from release_date
month_base = merged[["release_date"]].copy()
month_base["release_dt"] = pd.to_datetime(month_base["release_date"], errors="coerce")
month_base = month_base.dropna(subset=["release_dt"])
month_base["release_month"] = month_base["release_dt"].dt.month


# try first gross first
month_df = month_base.join(merged[gross_pref].rename("gross"))
month_df["gross"] = pd.to_numeric(month_df["gross"], errors="coerce")
month_df = month_df.dropna(subset=["gross"])
month_df = month_df[month_df["gross"] > 0]

months_present = month_df["release_month"].nunique()
used_col = gross_pref

# if data missing, fall back to budgets dataset gross
if months_present <= 2 and grossFallback is not None:
    month_df = month_base.join(merged[grossFallback].rename("gross"))
    month_df["gross"] = pd.to_numeric(month_df["gross"], errors="coerce")
    month_df = month_df.dropna(subset=["gross"])
    month_df = month_df[month_df["gross"] > 0]
    used_col = grossFallback

# compute monthly averages show all months & fill missing with 0
month_avg = (
    month_df.groupby("release_month")["gross"]
    .mean()
    .reindex(range(1, 13), fill_value=0)
)

plt.figure(figsize=(12, 6))
plt.bar(month_avg.index, month_avg.values)
plt.xlabel("Release Month")
plt.ylabel("Average Worldwide Revenue ($)")
plt.title(f"Average Worldwide Revenue by Release Month")
plt.gca().yaxis.set_major_formatter(FuncFormatter(dollars_millions))
plt.xticks(range(1, 13))
plt.tight_layout()
plt.savefig(out_dir + "avg_worldwide_revenue_by_month.png")
plt.show()
```



## 4. Conclusions & Key Findings

1. High ratings ≠ high revenue. Audience ratings and worldwide gross show only a weak relationship. This simply great reviews don't guarantee box office success. This is because commercial performance is driven more by distribution, marketing, franchise power, and release timing other than by ratings alone.
2. Certain genres consistently outperform financially. A small set of genres delivers the strongest returns on investment. These genres combine scalable audiences with predictable demand, making them lo
3. Release timing materially affects revenue. Average worldwide revenue varies significantly by release month. Peak months which are usually major holidays and summer windows outperform off-peak months greatly. This confirms that when a film is released can be as important as what film is released.
4. International markets drive the majority of revenue. Foreign gross often exceeds domestic gross for top-performing genres and titles. This indicates that global appeal and international distribution strategy are critical to maximizing returns. Also global reach may mean having the film scripted in english or atleast offer english subtitles.
5. ROI highlights capital efficiency, not just raw revenue. Big-budget blockbusters dominate absolute revenue, but ROI analysis shows that mid-budget films in the right genres can outperform in capital efficiency.

In [ ]: